

Transferência de Aprendizado em Deep Learning Utilizando Tensorflow Hub

IV MEETUP PYDATA MANAUS








FAGNER CUNHA



DEEP LEARNING




REDES NEURAIS COM MUITAS CAMADAS



Esquema da InceptionV3
Fonte: Adaptado de [1]

- Novas técnicas de treinamento
- Poder computacional
- **Muitos** dados





BASES DE DADOS

The logo for ImageNet, featuring the word "IMAGENET" in a sans-serif font. The letter "A" is replaced by a small graphic of three colored squares (green, orange, red) arranged in a triangular pattern.

1,2M IMAGENS



COCO
Common Objects in Context

330K IMAGENS
1,5M INSTÂNCIAS



WIKIPÉDIA
A enciclopédia livre

2500M PALAVRAS

BASES DE DADOS



1,2M IMAGENS




330K IMAGENS
1,5M INSTÂNCIAS



2500M PALAVRAS

Minha base é **pequena!**

TRANSFERÊNCIA DE APRENDIZADO

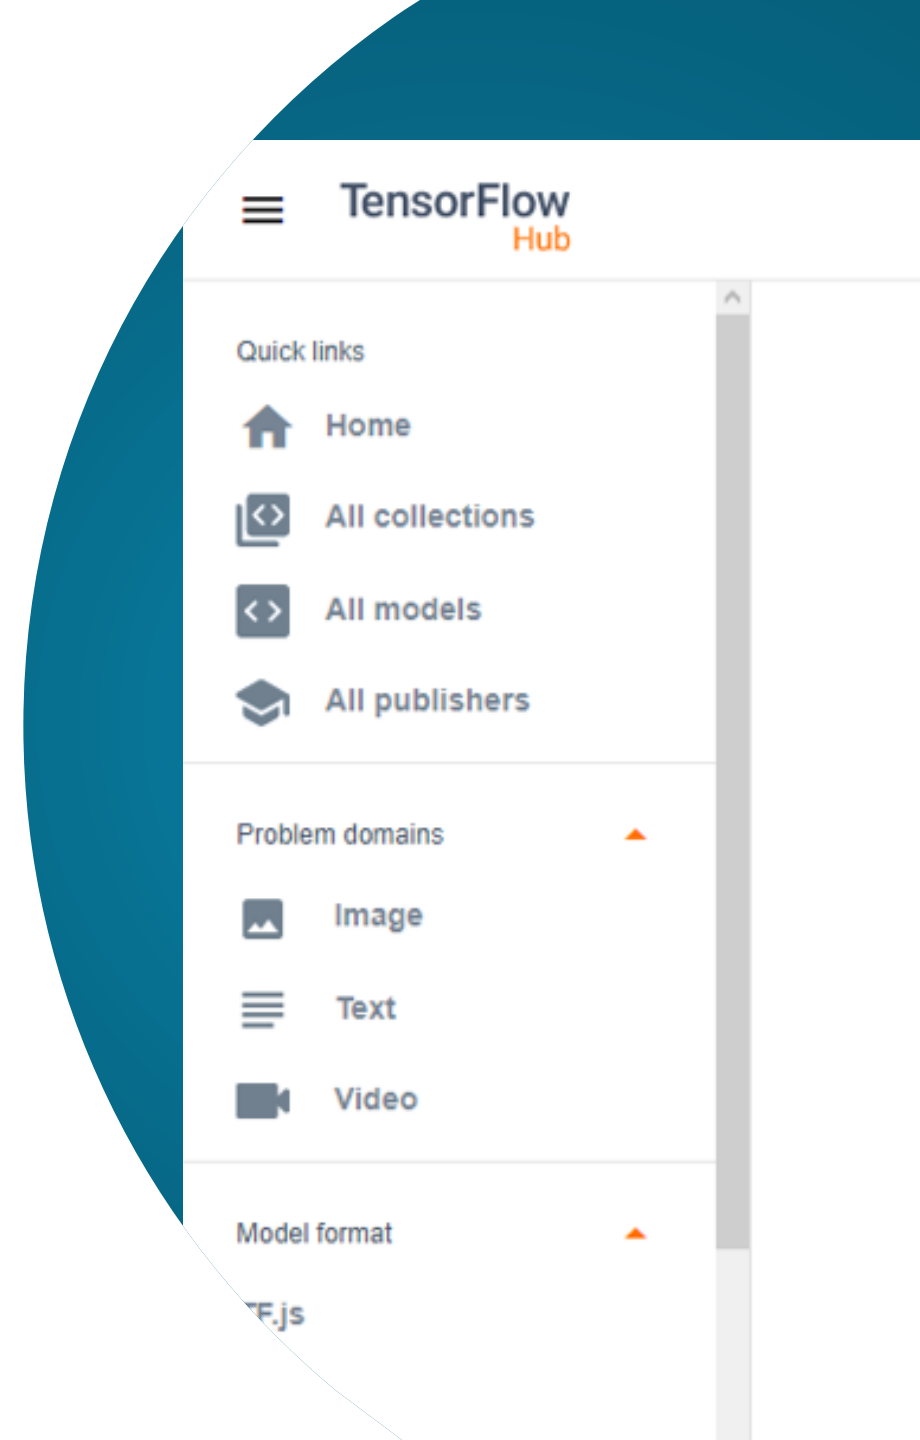





TENSORFLOW HUB

- Facilita o reuso de módulos pré-treinados
- Mais de 500 módulos pra problemas nos domínios de texto, imagem, vídeo e áudio

- tfhub.dev



TFHUB.DEV



mobilenet

Send feedback

Filters

Clear all

Problem domain

Model format

TF.js

TFLite

Coral

TF Version

TF1

TF2

Fine tunable

Architecture

Publisher

Dataset

MobileNet V2 | ImageNet (ILSVRC-2012-CLS)

MobileNet V2 | ImageNet (ILSVRC-2012-CLS)

Image feature vector

imagenet/mobilenet_v2_100_128...

Published by: Google Updated: 04/30/2020

Feature vectors of images with MobileNet V2 (depth multiplier 1.00) trained on ImageNet (ILSVRC-2012-CLS).

MobileNet V2 | ImageNet (ILSVRC-2012-CLS)

Image feature vector

imagenet/mobilenet_v2_100_224...

Published by: Google Updated: 05/01/2020

Feature vectors of images with MobileNet V2 (depth multiplier 1.00) trained on ImageNet (ILSVRC-2012-CLS).

MobileNet V2 | ImageNet (ILSVRC-2012-CLS)

Image classification

imagenet/mobilenet_v2_100_96...

Published by: Google Updated: 04/30/2020

Imagenet (ILSVRC-2012-CLS) classification with MobileNet V2 (depth multiplier 1.00).

MobileNet V2 | ImageNet (ILSVRC-2012-CLS)

Image classification


imagenet/mobilenet_v2_100_160...


Published by: Google Updated: 05/01/2020

Imagenet (ILSVRC-2012-CLS) classification with MobileNet V2 (depth multiplier 1.00).

MobileNet V2 | ImageNet (ILSVRC-2012-CLS)

TFHUB.DEV



 TensorFlow Hub

mobilenet

Send feedback

← imagenet/mobilenet_v2_100_224/feature_vector

Problem domain

Image feature vector

Architecture

MobileNet V2

Publisher


Google

Dataset

ImageNet (ILSVRC-2012-CLS)

Format: TF2.0 Saved Model

Fine tunable: Yes

License: [Apache-2.0](#) 

Model formats

TF2.0 Saved Model

.JS (v1, default/1)

.JS (v2, default/1)

.JS (v3, default/1)

Want to use this model?

To use this model, take a look at the example code, or at [our user guide](#).


You can also try out the associated Colab.

Copy URL to clipboard

Download Model

Open Colab Notebook

Asset size: 8.12MB



USANDO TFHUB

In [5]: `import tensorflow_hub as hub`

```
from tensorflow.keras.layers import Dense
from tensorflow.keras import Sequential
```

In [6]: `MODEL_URL = 'https://tfhub.dev/google/imagenet/mobilenet_v2_100_224/feature_vector/4'`
`feature_extractor = hub.KerasLayer(MODEL_URL,`
 `input_shape=(224, 224, 3),`
 `trainable=False)`

In [7]: `model = Sequential([`
 `feature_extractor,`
 `Dense(NUM_CLASSES)`
`])`
`model.summary()`


Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
keras_layer (KerasLayer)	(None, 1280)	2257984
dense (Dense)	(None, 3)	3843
=====		

Total params: 2,261,827


Trainable params: 3,843

Non-trainable params: 2,257,984



TENSORFLOW DATASETS

- Mais de 170 datasets (áudio, imagem, vídeo, texto)
- <https://www.tensorflow.org/datasets/catalog/>
- Fornecidos como **tf.data.Datasets**
- Dataset de classificação de imagens: `rock_paper_scissors [2]`





USANDO 0 TENSORFLOW DATASETS

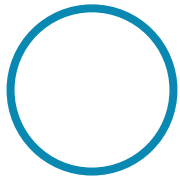


```
In [8]: import tensorflow_datasets as tfds
```

```
In [9]: (jokempo_train, jokempo_test), info = tfds.load('rock_paper_scissors',  
                                                    split=['train', 'test'],  
                                                    shuffle_files=True,  
                                                    as_supervised=True,  
                                                    with_info=True)
```

```
In [10]: info
```

```
Out[10]: tfds.core.DatasetInfo(  
    name='rock_paper_scissors',  
    version=3.0.0,  
    description='Images of hands playing rock, paper, scissor game.',  
    homepage='http://laurencemoroney.com/rock-paper-scissors-dataset',  
    features=FeaturesDict({  
        'image': Image(shape=(300, 300, 3), dtype=tf.uint8),  
        'label': ClassLabel(shape=(), dtype=tf.int64, num_classes=3),  
    }),  
    total_num_examples=2892,  
    splits={  
        'test': 372,  
        'train': 2520,  
    },  
    supervised_keys=('image', 'label'),
```



USANDO 0 TENSORFLOW DATASETS

```
In [11]: show_examples(jokempo_train, info)
```



rock (0)



scissors (2)



scissors (2)



paper (1)



paper (1)



paper (1)




rock (0)




paper (1)



TF.DATA



- API para construir pipelines de dados
- Pipelines complexos a partir de operações simples
- Flexibilidade
- Fornecer dados em tempo hábil para GPU





TF.DATA



```
In [13]: IMAGE_HEIGHT = 224
         IMAGE_WIDTH = 224
         IMAGE_PRE_CROP_WIDTH = 240
         IMAGE_PRE_CROP_HEIGHT = 240
         ROTATION_ANGLE = 45
         AUTOTUNE = tf.data.experimental.AUTOTUNE
```

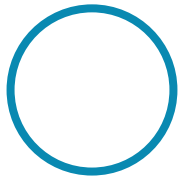
```
In [14]: def read_decode_image(image, label):
         image = tf.cast(image, tf.float32) / 255.

         return image, label
```

```
In [15]: def apply_augmentation(image, label):
         rotation_theta = deg2rad(ROTATION_ANGLE)


         image = tf.image.resize(image, [IMAGE_PRE_CROP_WIDTH, IMAGE_PRE_CROP_HEIGHT])
         image = tf.image.random_flip_left_right(image)
         image = tfa.image.rotate(image,
                                   tf.random.uniform(shape=[1],
                                                       minval=-rotation_theta,
                                                       maxval=rotation_theta),
                                   interpolation='BILINEAR')
         image = tf.image.random_crop(image, size=[IMAGE_WIDTH, IMAGE_HEIGHT, 3])

         return image, label
```





TF.DATA



```
In [17]: show_examples(jokempo_train, info)
```




rock (0)



rock (0)



paper (1)




paper (1)




paper (1)




scissors (2)




rock (0)



paper (1)






TF.DATA

```
In [18]: TRAIN_LENGTH = info.splits['train'].num_examples  
TEST_LENGTH = info.splits['test'].num_examples  
NUM_CLASSES = info.features['label'].num_classes
```

```
In [19]: BATCH_SIZE = 32  
EPOCHS = 10  
AUTOTUNE = tf.data.experimental.AUTOTUNE
```


Pipeline de treinamento:

```
In [20]: jokempo_train = jokempo_train.map(read_decode_image, num_parallel_calls=AUTOTUNE)  
jokempo_train = jokempo_train.map(apply_augmentation, num_parallel_calls=AUTOTUNE)  
jokempo_train = jokempo_train.repeat(EPOCHS)  
jokempo_train = jokempo_train.shuffle(buffer_size=TRAIN_LENGTH)  
jokempo_train = jokempo_train.batch(BATCH_SIZE, drop_remainder=True)  
jokempo_train = jokempo_train.prefetch(buffer_size=AUTOTUNE)
```






TF.DATA



Pipeline de teste:

```
In [22]: def resize_image(image, label):  
         image = tf.image.resize(image, [IMAGE_WIDTH, IMAGE_HEIGHT])  
  
         return image, label
```

```
In [23]: jokempo_test = jokempo_test.map(read_decode_image, num_parallel_calls=AUTOTUNE)  
jokempo_test = jokempo_test.map(resize_image, num_parallel_calls=AUTOTUNE)  
jokempo_test = jokempo_test.repeat(EPOCHS)  
jokempo_test = jokempo_test.batch(BATCH_SIZE, drop_remainder=True)  
jokempo_test = jokempo_test.prefetch(buffer_size=AUTOTUNE)
```





TREINANDO O MODELO



```
In [21]: history = model.fit(jokempo_train,
                             steps_per_epoch=TRAIN_LENGTH // BATCH_SIZE,
                             epochs=EPOCHS,
                             validation_data=jokempo_test,
                             validation_steps=TEST_LENGTH // BATCH_SIZE,
                             initial_epoch=0)
```

Train for 78 steps, validate for 11 steps

Epoch 1/10

78/78 [=====] - 10s 134ms/step - loss: 0.4306 - accuracy: 0.8598 - val_loss: 0.2708 - val_accuracy: 0.9403

Epoch 2/10

78/78 [=====] - 4s 49ms/step - loss: 0.1066 - accuracy: 0.9872 - val_loss: 0.3014 - val_accuracy: 0.9119

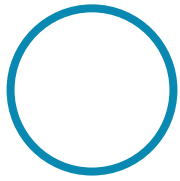
Epoch 3/10

78/78 [=====] - 4s 51ms/step - loss: 0.0642 - accuracy: 0.9912 - val_loss: 0.2601 - val_accuracy: 0.9233

Epoch 4/10

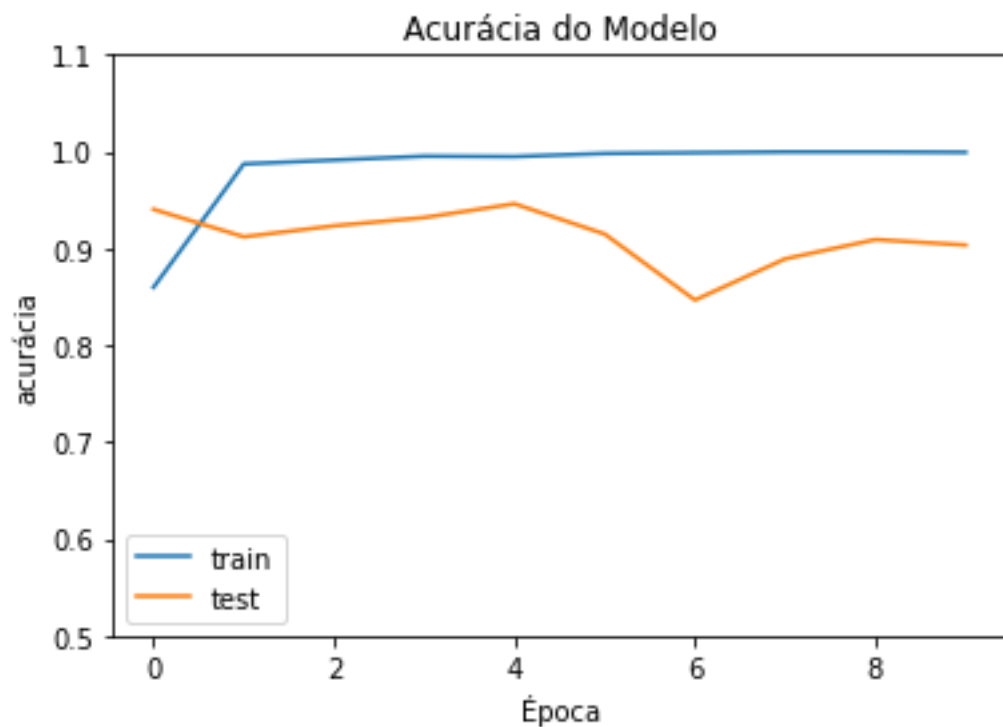
78/78 [=====] - 4s 51ms/step - loss: 0.0515 - accuracy: 0.9952 - val_loss: 0.2212 - val_accuracy: 0.9318

Epoch 5/10

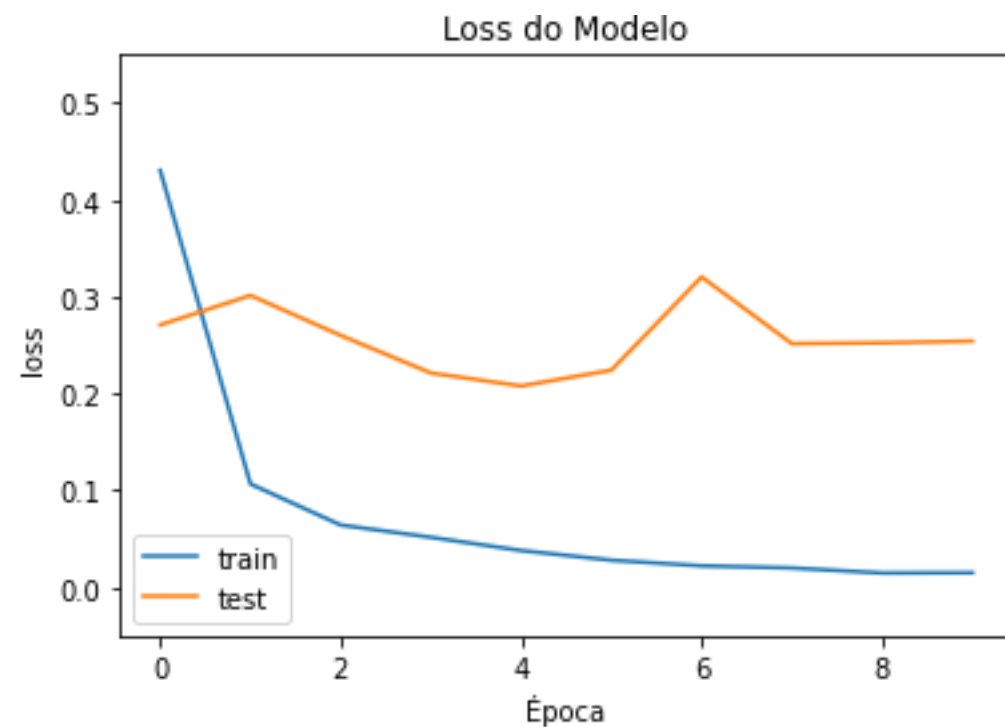



VISUALIZANDO OS RESULTADOS

```
In [22]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.ylim([0.5, 1.1])
plt.title('Acurácia do Modelo')
plt.ylabel('acurácia')
plt.xlabel('Época')
plt.legend(['train', 'test'], loc='lower left')
plt.show()
```




```
In [23]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.ylim([-0.05, 0.55])
plt.title('Loss do Modelo')
plt.ylabel('loss')
plt.xlabel('Época')
plt.legend(['train', 'test'], loc='lower left')
plt.show()
```





SÓ MUDAR A ÚLTIMA CAMADA?

- Adicionar camadas intermediárias após o extrator de características
- Fine tuning
- Utilizar como gerador de features para outros modelos como o SVM
- Treinar modelos de deep learning é quase arte





OBRIGADO !



LinkedIn:

<https://www.linkedin.com/in/cunhaf/>



Twitter

@fagner_cunha




E-mail

fagner.cunha@icomp.ufam.edu.br




Github

@alcunha



REFERÊNCIAS

- [1] Inception in TensorFlow. Disponível em: <https://github.com/tensorflow/models/tree/master/research/inception>. Acesso em: 30 de abril de 2020.
- [2] Moroney, Laurence (2019). Rock, Paper, Scissors Dataset. Disponível em: <http://laurencemoroney.com/rock-paper-scissors-dataset>. Acesso em: 30 de abril de 2020.
- [3] Transfer learning with TensorFlow Hub. Disponível em: https://www.tensorflow.org/tutorials/images/transfer_learning_with_hub. Acesso em: 30 de abril de 2020.





PERGUNTAS

